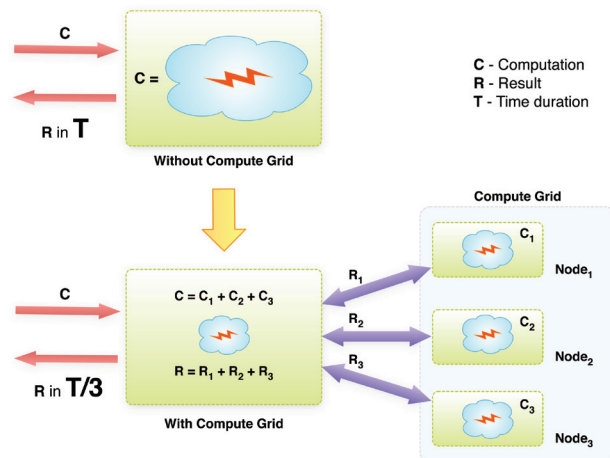


GridGain + Big Data = Fast Data

GridGain: In-Memory Compute and Data Grid Technologies

Compute Grid:

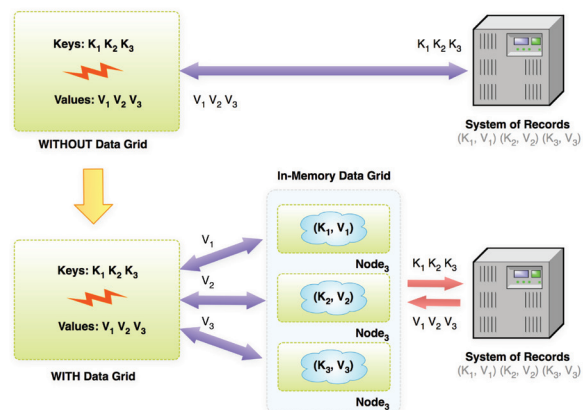
- Direct API support for MapReduce
- Distributed closure/function execution
- Pluggable failover resolution
- Pluggable topology resolution
- Distributed task session
- Annotation-based execution
- Asynchronous execution
- Redundant job mapping
- Continuous job mapping
- Full and partial asynchronous aggregation
- Adaptive and dynamic task split
- Checkpoints for long-running tasks
- Early and late load balancing
- Affinity co-location with data grids
- AOP-based cloud enabling
- Pluggable early and late load balancing



- Customizable failover resolution
- Redundant jobs
- Asynchronous results processing

In-Memory Data Grid:

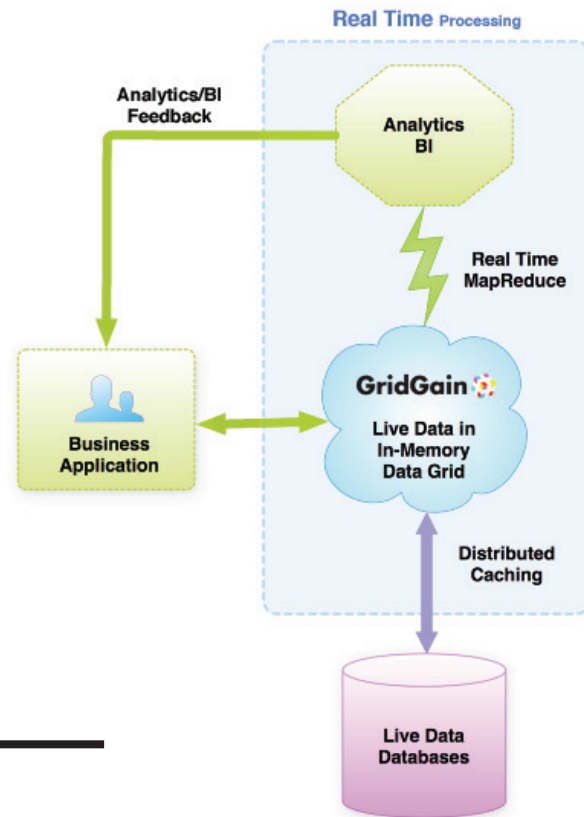
- Local caches
- Fully replicated caches
- Partitioned caches based on consistent hashing
- Data replication and invalidation
- Sync and Async replication and invalidation
- Data backups for partitioned caches
- Pessimistic distributed locks
- Pessimistic transactions
- Optimistic transactions
- Eventually consistent transactions
- Isolation levels:
 - Read_Committed, Repeatable_Read, Serializable
- JTA/JCA integration
- Expiration policies (LRU, LIRS, FIFO, Random, Time)
- Named caches (multiple caches coexisting)
- Read-Through and Write-Through behavior
- Synchronous and asynchronous APIs
- Cached data utilizes Zero Deployment



- Pluggable data overflow storage
- Write-From-Behind support
- Support for indexed unstructured data
- Distributed Queries: SQL, Lucene, H2 and text based
- Local and remote filtering, transformation and reduction
- Paginated query result sets
- Pluggable topology segmentation handling built-in distributed data structures

Using GridGain your Real Time Big Data applications can:

- Work in a zero-deployment mode
- Easily scale up or down based on demand
- Cache distributed data in an In-Memory data grid
- Co-locate data and computations
- Run SQL queries against cached data
- Store and query JSON objects
- Speed up processing using real-time map/reduce
- Use distributed thread pools
- Distribute the workload on the grid
- Effectively exchange messages
- Auto-discover all grid resources
- Execute closures on the grid
- Grid-enable Java, Groovy and Scala



What the users say:


"Overall we had an awesome experience with GridGain. The grid application ran flawlessly during our busiest time of the year (Thanksgiving weekend). It ran so flawlessly, that this morning I forgot which boxes it was physically running on."

"I was able to watch a few of your videos (installing GridGain and an application in 15 minutes) and had the hello world example running in no time. You and your team have done a great job making GridGain accessible, well documented and easy to use."

"Forget IBM, build your own Blue Gene for under \$50,000:-)."

"...just a quick note to tell you guys that my first experiences with GridGain are absolutely great. I implemented a prototype of one of our use cases (validating millions of image links) in about a day and had got it up and running on a 16 node EC2 cluster very easily."

"I got a working instance of a prototype hooked up to my front end in a matter of an hour. And it just works. Five stars guys and keep it up :-)."



Contact Us:

info@gridgain.com ■ @gridgain
ph 650-241-2281 ■ fax 925-369-7193

GridGain.com
1065 East Hillsdale Blvd. ■ Suite 230
Foster City, CA ■ 94404