

# Postgres Demystified



Craig Kerstiens

@craigkerstiens


<http://www.craigkerstiens.com>

<https://speakerdeck.com/u/craigkerstiens/p/postgres-demystified>

# Postgres Demystified



# Postgres Demystified



 heroku**postgres**

# Getting Setup



Postgres.app

# Agenda

Brief History

Developing w/ Postgres

Postgres Performance

Querying

# Postgres History

Postgres  
PostgreSQL

Post Ingress  
Around since 1989/1995  
Community Driven/Owned

# MVCC

Each query sees transactions committed before it  
Locks for writing don't conflict with reading

# Why Postgres



# Why Postgres

“its the emacs of databases”

# Developing w/ Postgres

# Basics

*psql is your friend*

# Datatypes

A word cloud of database datatypes in teal color, arranged in a circular pattern around the center. The datatypes include: timestamp, array, date, UUID, boolean, integer, interval, XML, enum, char, float, circle, point, numeric, money, polygon, bytea, serial, line, smallint, bigint, inet, cidr, path, macaddr, time, text, tsquery, timestamp, timetz, box, tsvector, and varchar.

timestamp array date UUID boolean integer interval XML enum char float circle point numeric money polygon bytea serial line smallint bigint inet cidr path macaddr time text tsquery timestamp timetz box tsvector varchar

# Datatypes

timestampz array date interval UUID boolean  
smallint line bigint XML enum integer  
serial bytea money char  
inet cidr varchar polygon numeric point  
path macaddr time text tsquery float  
macaddr time text tsquery timestamp timetz circle  
tsvector

# Datatypes

UUID  
boolean  
integer  
interval  
date  
bigint  
array  
timestampz  
smallint  
XML  
enum  
char  
float  
circle  
point  
money  
numeric  
polygon  
text  
tsquery  
timestamp  
timetz  
box  
time  
tsvector  
varchar  
bytea  
serial  
line  
inet  
cidr  
macaddr  
path

# Datatypes

timestamp, array, date, UUID, boolean, integer, interval, XML, enum, bigint, money, char, float, circle, point, polygon, line, path, box, tsvector, time, text, tsquery, timestamp, timetz, cidr, macaddr, inet, serial, bytea, numeric, varchar, smallint, timestampz

# Datatypes

```
CREATE TABLE items (  
  id serial NOT NULL,  
  name varchar (255),  
  tags varchar(255) [],  
  created_at timestamp  
);
```



# Datatypes

```
CREATE TABLE items (  
  id serial NOT NULL,  
  name varchar (255),  
  tags varchar(255) [],  
  created_at timestamp  
);
```

# Datatypes

```
CREATE TABLE items (  
  id serial NOT NULL,  
  name varchar (255),  
  tags varchar(255) [],  
  created_at timestamp  
);
```

# Datatypes

```
INSERT INTO items  
VALUES (1, 'Ruby Gem', '{"Programming","Jewelry"}', now());
```

```
INSERT INTO items  
VALUES (2, 'Django Pony', '{"Programming","Animal"}', now());
```

# Datatypes

timestamptz array date UUID boolean  
interval integer  
bigint XML enum  
smallint line money char  
serial bytea point float  
inet cidr varchar polygon numeric circle  
path macaddr time text tsquery timestamp  
tsvector box

# Datatypes wish list

email

url

zip

phone

# Extensions

A word cloud of PostgreSQL extensions. The words are arranged in a scattered, non-uniform pattern across the slide. The extensions listed are: dblink, citext, isn, cube, unaccent, hstore, ltree, tablefunc, uuid-oss, pgcrypto, earthdistance, btree\_gist, trigram, pgrowlocks, fuzzystmatch, dict\_int, dict\_xsyn, and pgstattuple. The words are in a dark teal color and vary in size, with 'Extensions' being the largest at the top.

dblink hstore uuid-oss trigram pgstattuple  
citext pgrowlocks  
isn pgcrypto fuzzystmatch  
cube ltree earthdistance dict\_int  
unaccent tablefunc dict\_xsyn btree\_gist

# Extensions

dblink hstore uuid-oss trigram pgstattuple  
citext pgcrypto pgrowlocks  
isn ltree fuzzystmatch  
cube earthdistance dict\_int  
unaccent tablefunc dict\_xsyn  
btree\_gist

# NoSQL in your SQL

```
CREATE EXTENSION hstore;  
CREATE TABLE users (  
    id integer NOT NULL,  
    email character varying(255),  
    data hstore,  
    created_at timestamp without time zone,  
    last_login timestamp without time zone  
);
```



# hStore

```
INSERT INTO users  
VALUES (  
  1,  
  'craig.kerstiens@gmail.com',  
  'sex => "M", state => "California",  
  now(),  
  now()  
);
```

# JSON

SELECT

```
'{"id":1,"email": "craig.kerstiens@gmail.com",}'::json;
```

9.2

# JSON

SELECT

```
'{"id":1,"email": "craig.kerstiens@gmail.com",}'::json;
```

## V8 w/ PLV8

9.2

# JSON

SELECT

```
'{"id":1,"email": "craig.kerstiens@gmail.com",}'::json;
```

## V8 w/ PLV8

Bad Idea

9.2

# JSON

SELECT

```
'{"id":1,"email": "craig.kerstiens@gmail.com",}'::json;
```

## V8 w/ PLV8

```
create or replace function  
js(src text) returns text as $$  
  return eval(  
    "(function() { " + src + "})"  
  );  
$$ LANGUAGE plv8;
```

Bad Idea

9.2

# Range Types

**9.2**

# Range Types

```
CREATE TABLE talks (room int, during tsrange);  
INSERT INTO talks VALUES  
  (3, '[2012-09-24 13:00, 2012-09-24 13:50)');
```

# Range Types

```
CREATE TABLE talks (room int, during tsrange);  
INSERT INTO talks VALUES  
  (3, '[2012-09-24 13:00, 2012-09-24 13:50)');
```

```
ALTER TABLE talks ADD EXCLUDE USING gist (during WITH &&);  
INSERT INTO talks VALUES  
  (1108, '[2012-09-24 13:30, 2012-09-24 14:00)');  
ERROR: conflicting key value violates exclusion constraint  
"talks_during_excl"
```



# Full Text Search

# Full Text Search

TSVECTOR - Text Data

TSQUERY - Search Predicates

Specialized Indexes and Operators

# Datatypes

PostgreSQL data types word cloud:

- uuid
- boolean
- integer
- enum
- char
- float
- circle
- box
- timestamp
- tsvector
- time
- text
- tsquery
- numeric
- point
- money
- bytea
- serial
- inet
- path
- macaddr
- cidr
- varchar
- poligon
- line
- smallint
- timestampz
- array
- bigint
- date
- interval
- XML

# PostGIS

# PostGIS

I. New datatypes      i.e. (2d/3d boxes)

# PostGIS

1. New datatypes      i.e. (2d/3d boxes)
2. New operators      i.e. `SELECT foo && bar ...`

# PostGIS

1. New datatypes *i.e. (2d/3d boxes)*
2. New operators *i.e. SELECT foo && bar ...*
3. Understand relationships and distance  
*i.e. person within location, nearest distance*





# Performance

# Sequential Scans

# Sequential Scans

They're Bad

# Sequential Scans

They're Bad (most of the time)

# Indexes

# Indexes

They're Good

# Indexes

They're Good (most of the time)

# Indexes

B-Tree

Generalized Inverted Index (GIN)

Generalized Search Tree (GIST)

K Nearest Neighbors (KNN)

Space Partitioned GIST (SP-GIST)



# Indexes

## B-Tree

Default

Usually want this

# Indexes

## Generalized Inverted Index (GIN)

Use with multiple values in 1 column  
Array/hStore

# Indexes

## Generalized Search Tree (GIST)

Full text search

Shapes

# Understanding Query Perf

Given

```
SELECT last_name  
FROM employees  
WHERE salary >= 50000;
```

# Explain

```
# EXPLAIN SELECT last_name FROM employees WHERE salary >= 50000;
```

## QUERY PLAN

-----  
Seq Scan on employees (cost=0.00..35811.00 rows=1 width=6)  
 Filter: (salary >= 50000)  
(3 rows)

# Explain

```
# EXPLAIN SELECT last_name FROM employees WHERE salary >= 50000;
```

**Startup Cost**

QUERY PLAN

Seq Scan on employees (cost=0.00..35811.00 rows=1 width=6)  
Filter: (salary >= 50000)  
(3 rows)

# Explain

```
# EXPLAIN SELECT last_name FROM employees WHERE salary >= 50000;
```

**Startup Cost**

QUERY PLAN

**Max Time**

Seq Scan on employees (cost=0.00. 3581 1.00 rows=1 width=6)  
Filter: (salary >= 50000)  
(3 rows)

# Explain

```
# EXPLAIN SELECT last_name FROM employees WHERE salary >= 50000;
```

**Startup Cost**

QUERY PLAN

**Max Time**

Seq Scan on employees (cost=0.00.35811.00 rows=1 width=6)  
Filter: (salary >= 50000)  
(3 rows)

**Rows Returned**



# Explain

```
# EXPLAIN ANALYZE SELECT last_name FROM employees WHERE  
salary >= 50000;
```

## QUERY PLAN

-----  
Seq Scan on employees (cost=0.00..35811.00 rows=1 width=6) (actual  
time=2.401..295.247 rows=1428 loops=1)

Filter: (salary >= 50000)

Total runtime: 295.379

(3 rows)

# Explain

```
# EXPLAIN ANALYZE SELECT last_name FROM employees WHERE  
salary >= 50000;
```

**Startup Cost**

QUERY PLAN

Seq Scan on employees (cost=0.00..35811.00 rows=1 width=6) (actual  
time=**2.401**..295.247 rows=1428 loops=1)

Filter: (salary >= 50000)

Total runtime: 295.379

(3 rows)

# Explain

```
# EXPLAIN ANALYZE SELECT last_name FROM employees WHERE  
salary >= 50000;
```

**Startup Cost**

**Max Time**

Seq Scan on employees (cost=0.00..35811.00 rows=1 width=6) (actual  
time=2.401, 295.247 rows=1428 loops=1)

Filter: (salary >= 50000)

Total runtime: 295.379

(3 rows)

# Explain

```
# EXPLAIN ANALYZE SELECT last_name FROM employees WHERE  
salary >= 50000;
```

**Startup Cost**

**Max Time**

Seq Scan on employees (cost=0.00..35811.00 rows=1 width=6) (actual  
time=2.401, 295.247 rows=1428 loops=1)

Filter: (salary >= 50000)

Total runtime: 295.379  
(3 rows)

**Rows Returned**

# Explain

```
# EXPLAIN ANALYZE SELECT last_name FROM employees WHERE  
salary >= 50000;
```

**Startup Cost**

**Max Time**

Seq Scan on employees (cost=0.00..35811.00 rows=1 width=6) (actual  
time=2.401, 295.247 rows=1428 loops=1)

Filter: (salary >= 50000)

Total runtime: 295.379  
(3 rows)

**Rows Returned**

```
# CREATE INDEX idx_emps ON employees (salary);
```

```
# CREATE INDEX idx_emps ON employees (salary);  
# EXPLAIN ANALYZE SELECT last_name FROM employees WHERE  
salary >= 50000;
```

## QUERY PLAN

-----  
Index Scan using idx\_emps on employees (cost=0.00..8.49 rows=1  
width=6) (actual time = 0.047..1.603 rows=1428 loops=1)  
  Index Cond: (salary >= 50000)  
Total runtime: 1.771 ms  
(3 rows)

```
# CREATE INDEX idx_emps ON employees (salary);  
# EXPLAIN ANALYZE SELECT last_name FROM employees WHERE  
salary >= 50000;
```

## QUERY PLAN

-----  
Index Scan using idx\_emps on employees (cost=0.00..8.49 rows=1  
width=6) (actual time = 0.047..1.603 rows=1428 loops=1)

Index Cond: (salary >= 50000)

Total runtime: 1.771 ms  
(3 rows)



# Indexes Pro Tips

# Indexes Pro Tips

**CREATE INDEX CONCURRENTLY**

# Indexes Pro Tips

CREATE INDEX CONCURRENTLY

CREATE INDEX WHERE foo=bar

# Indexes Pro Tips

CREATE INDEX CONCURRENTLY

CREATE INDEX WHERE foo=bar

SELECT \* WHERE foo LIKE '%bar%' is BAD

# Indexes Pro Tips

CREATE INDEX CONCURRENTLY


CREATE INDEX WHERE foo=bar

SELECT \* WHERE foo LIKE '%bar%' is BAD

SELECT \* WHERE Food LIKE 'bar%' is OKAY

# Extensions

dblink hstore uuid-ossop trigram pgstattuple  
citext pgcrypto pgrowlocks  
isn ltree fuzzystmatch  
cube earthdistance dict\_int  
unaccent tablefunc btree\_gist dict\_xsyn



# Index Hit Rate

```
SELECT
  relname,
  100 * idx_scan / (seq_scan + idx_scan),
  n_live_tup
FROM pg_stat_user_tables
ORDER BY n_live_tup DESC;
```



# Cache Hit Rate

```
SELECT
  relname::text,
  heap_blks_read + heap_blks_hit as reads,
  round(100 * heap_blks_hit / (heap_blks_hit + heap_blks_read)) as
  hit_pct,
  round(100 * idx_blks_hit / (idx_blks_hit + idx_blks_read)) as idx_hit_pct
FROM   pg_statio_user_tables
WHERE  heap_blks_hit + heap_blks_read + idx_blks_hit + idx_blks_read
> 0
ORDER BY 2 DESC;
```

# pg\_stats\_statements

**9.2**

# pg\_stats\_statements

```
$ select * from pg_stat_statements where query ~ 'from users where email';
```

userid		16384
dbid		16388
query		select * from users where email = ?;
calls		2
total_time		0.000268
rows		2
shared_blks_hit		16
shared_blks_read		0
shared_blks_dirtied		0
shared_blks_written		0
local_blks_hit		0
local_blks_read		0
local_blks_dirtied		0
local_blks_written		0
temp_blks_read		0
temp_blks_written		0
time_read		0
time_write		0

9.2

# pg\_stats\_statements

**9.2**

# pg\_stats\_statements

```
SELECT query, calls, total_time, rows, 100.0 * shared_blks_hit /  
       nullif(shared_blks_hit + shared_blks_read, 0) AS hit_percent  
FROM pg_stat_statements ORDER BY total_time DESC LIMIT 5;
```

-----  
query | UPDATE pgbench\_branches SET bbalance = bbalance + ?  
WHERE bid = ?;

calls | 3000

total\_time | 9609.001000000002

rows | 2836

hit\_percent | 99.9778970000200936

9.2



# Querying

# Window Functions

Example:

Biggest spender by state



# Window Functions

```
SELECT
  email,
  users.data->'state',
  sum(total(items)),
  rank() OVER
    (PARTITION BY users.data->'state'
     ORDER BY sum(total(items)) desc)
FROM
  users, purchases
WHERE purchases.user_id = users.id
GROUP BY 1, 2;
```

# Window Functions

```
SELECT
  email,
  users.data->'state',
  sum(total(items)),
  rank() OVER
    (PARTITION BY users.data->'state'
     ORDER BY sum(total(items)) desc)
FROM
  users, purchases
WHERE purchases.user_id = users.id
GROUP BY 1, 2;
```

# Extensions

dblink hstore uuid-oss trigram pgstattuple  
citext pgrowlocks  
isn pgcrypto fuzzystrmatch  
ltree  
cube earthdistance dict\_int  
unaccent tablefunc btree\_gist dict\_xsyn

# Fuzzystmatch

# Fuzzystmatch

```
SELECT soundex('Craig'), soundex('Will'), difference('Craig', 'Will');
```

# Fuzzystmatch

```
SELECT soundex('Craig'), soundex('Will'), difference('Craig', 'Will');
```

```
SELECT soundex('Craig'), soundex('Greg'), difference('Craig', 'Greg');  
SELECT soundex('Willl'), soundex('Will'), difference('Willl', 'Will');
```

# Moving Data Around

```
\copy (SELECT * FROM users) TO '~/users.csv';
```

```
\copy users FROM '~/users.csv';
```

# db\_link

```
SELECT dblink_connect('myconn', 'dbname=postgres');  
SELECT * FROM dblink('myconn', 'SELECT * FROM foo') AS t(a int, b text);
```

a		b
1		example
2		example2



# Foreign Data Wrappers

oracle

mysql

odbc

twitter

sybase

redis

jdbc

files

couch

s3

www

ldap

informix

# Foreign Data Wrappers

```
CREATE EXTENSION redis_fdw;
```

```
CREATE SERVER redis_server  
  FOREIGN DATA WRAPPER redis_fdw  
  OPTIONS (address '127.0.0.1', port '6379');
```

```
CREATE FOREIGN TABLE redis_db0 (key text, value text)  
  SERVER redis_server  
  OPTIONS (database '0');
```

```
CREATE USER MAPPING FOR PUBLIC  
  SERVER redis_server  
  OPTIONS (password 'secret');
```

# Query Redis from Postgres

```
SELECT *  
FROM redis_db0;
```

```
SELECT  
  id,  
  email,  
  value as visits  
FROM  
  users,  
  redis_db0  
WHERE ('user_' || cast(id as text)) = cast(redis_db0.key as text)  
      AND cast(value as int) > 10;
```

# Readability

# Readability

```
WITH top_5_products AS (  
  SELECT products.*, count(*)  
  FROM products, line_items  
  WHERE products.id = line_items.product_id  
  GROUP BY products.id  
  ORDER BY count(*) DESC  
  LIMIT 5  
)  
  
SELECT users.email, count(*)  
FROM users, line_items, top_5_products  
WHERE line_items.user_id = users.id  
  AND line_items.product_id = top_5_products.id  
GROUP BY I  
ORDER BY I;
```

# Common Table Expressions

```
WITH top_5_products AS (  
  SELECT products.*, count(*)  
  FROM products, line_items  
  WHERE products.id = line_items.product_id  
  GROUP BY products.id  
  ORDER BY count(*) DESC  
  LIMIT 5  
)  
  
SELECT users.email, count(*)  
FROM users, line_items, top_5_products  
WHERE line_items.user_id = users.id  
  AND line_items.product_id = top_5_products.id  
GROUP BY I  
ORDER BY I;
```

Brief History  
Developing w/ Postgres  
Postgres Performance  
Querying

# Extras



# Extras

Listen/Notify

# Extras

Listen/Notify

Per Transaction Synchronous Replication

# Extras

Listen/Notify

Per Transaction Synchronous Replication

Drop index concurrently

# Postgres - TLDR

Datatypes

Conditional Indexes

Transactional DDL

Foreign Data Wrappers

Concurrent Index Creation

Extensions

Common Table Expressions

Fast Column Addition

Listen/Notify

Table Inheritance

Per Transaction sync replication

Window functions

NoSQL inside SQL

Momentum